



1. Masked average pooling: one embedding per neuron, per frame

encoder feature map (32 channels per location)

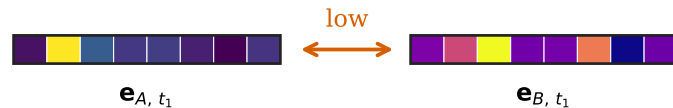


2. Compare embeddings by cosine similarity

Within-neuron (same neuron, different frames)



Between-neuron (different neurons, same frame)



gap = within - between (larger $\Rightarrow$ more stable, more distinct)